

**ACDA 2025 Mini-tutorial on Learning-Augmented Algorithms**

**Part 2**

# **Algorithmic Ideas from Learning-Augmented Algorithms**

**Vaidehi Srinivas**

Northwestern University

# Recap

## Learning-Augmented Algorithms (a.k.a. Algorithms with Predictions):

Take advantage of **unreliable predictions**, without sacrificing worst-case guarantees

### Consistency

If prediction is good, take advantage!

### Graceful Degradation

in magnitude of prediction error

### Robustness

If prediction is bad, revert to worst-case guarantee

## Ok great, but...

- Can we actually achieve this for interesting problems?
- Does it help develop **new** algorithmic ideas?

**Yes and yes!**

## Part 2

How to utilize extra **unreliable** information in algorithm design?

**Worst-case analysis**

Pessimistically ignore  
anything unreliable

**Heuristic**

Optimistically assume  
information is good, because  
it often is

**Learning-Augmented Algorithms**

Do both simultaneously!

### Goal of talk:

- Highlight a few interesting examples where **new theoretical frameworks** and abstractions lead to **new algorithmic ideas**
- Focus on data structures and optimizing runtime (though there's lots of exciting work in other relevant areas, like streaming, online algorithms, and more!)
- Exciting time to get involved!

# Three ideas

- (1) **Repeated Computations:** Sequences of related instances of a problem can be solved faster than one at a time
- (2) **Dynamic Algorithms/Data Structures:** Dynamic problems are easier with information about future updates
- (3) **Randomized Algorithms:** Randomized algorithms and data structures can be hedged to take advantage of extra information by incorporating a prior

# Three ideas

- (1) **Repeated Computations:** Sequences of related instances of a problem can be solved faster than one at a time
- (2) **Dynamic Algorithms/Data Structures:** Dynamic problems are easier with information about future updates
- (3) **Randomized Algorithms:** Randomized algorithms and data structures can be hedged to take advantage of extra information by incorporating a prior

# Sequences of Instances

**Setting:** solve a **sequence** of instances  $I_1, \dots, I_T$  of some algorithmic problem



**Example:** max-flow

**Goal:** minimize the **total runtime** to solve all of them

**Challenge:** If could solve sequences asymptotically faster than one at a time, would get an asymptotically faster worst-case algorithm

**Hope:** Adapt to **structure** in the sequence, when it exists!

(like dynamic algorithms, but want to impose less structure)


**Example:** suppose flow instances came from a predictable traffic network

**Formalism:** Learning-Augmented “**warm starts**” help us design and reason about algorithms for sequences of instances

# Warm Start Example: Max-Flow via Ford-Fulkerson

## Recall Ford-Fulkerson:

- Start with the all-0 flow
- Each iteration: find **augmenting path** to increase flow in residual graph by 1, in time  $O(|E|)$
- # of iterations: bounded by max flow in graph  $f$



**Total cost:**  $O(|E|f)$

## Warm start algorithm [Davies Moseley Vassilvitskii Wang '23]

- Start with a potentially **infeasible prediction**  $p$  for the flow on every edge
- Each iteration: augmenting-path like procedure in time  $O(|E|)$
- # of iterations: bounded by  $\eta = O(\sum_{e \in E} |p(e) - f(e)|)$ , where  $f(e)$  is the true flow on every edge ( $\ell_1$  error)

**Total cost:**  $O(|E| \cdot \min\{\eta, f\})$

**Always outputs correct answer!**

**Can we optimize runtime?**

# What to do with a warm start algorithm?

**Challenge:** a warm-start algorithm is only as good as the prediction!

For new instance, must predict the optimal flow

→ very high-dimensional learning problem

**Simple idea:** For a **sequence** of related instances, use  
“yesterday’s solution as today’s prediction”

[DMVW '23] Experiments on image segmentation tasks from frames of a video



(a) Image 1



(b) Image 5



(c) Image 10

Figure 4: Cuts (red) on the first, fifth, and last images from the  $120 \times 120$  pixels BIRDHOUSE sequence.



# Theoretical Interpretation

Is “yesterday’s solution as today’s prediction” theoretically principled?

Ford-Fulkerson Warm Start algorithm


[DMVW '23]:

- Flow solution is a vector indexed by edges
- Runtime of algorithm proportional to  $\ell_1$  **distance** between the predicted flow and the true solution

**Meta problem:** On each day  $t$ :

- Algorithm predicts a point  $P_t$  in the solution space
- True solution  $S_t$  is revealed
- Algorithm pays  $d(P_t, S_t)$

**Unlocks online algorithms toolkit!**




# Theoretical Interpretation

Is “yesterday’s solution as today’s prediction” theoretically principled?

- “Yesterday’s solution as today’s prediction” is competitive with the best fixed prediction in hindsight! [Khodak Balcan Talwalkar Vassilvitskii ’22]
  - Natural baseline from online algorithms literature
  - When solutions in sequence form a tight **cluster**, faster than solving instances one at a time

Can we do even better?

- Yes! Can design algorithms to take advantage of other forms of structure (multiple clusters) and compete against adaptive baselines [Blum Srinivas ’25]



# It's not just max-flow!

**Warm start algorithms** (with Learning-Augmented formalism) studied for:

- Ford-Fulkerson [Polak Zub '22][Davies Moseley Vassilvitskii Wang '23]
- Bipartite matching [Dinitz Im Lavastida Moseley Vassilvitskii '21][Chen Silwal Vakilian Zhang '22]
- Max-Flow via Push-Relabel [Davies Vassilvitskii Wang '24]

Can hope to prove for other **local search** algorithms, main new contribution is dealing with infeasibility

**In practice**, warm starts are used for many optimization problems!

**Takeaways:** Theoretical formalism of Learning-Augmented algorithms allows us to

- Analyze and compare strategies for solving sequences of related instances
- See new algorithmic opportunities
- **Modularize** (warm start vs. meta problem)

# References

**(1) Repeated Computations:** Sequences of related instances of a problem can be solved faster than one at a time

## **Learning-augmented warm start algorithms:**

- [Davies Moseley Vassilvitskii Wang '23][Polak Zub '22] Max-flow via Ford-Fulkerson
- [Dinitz Im Lavastida Moseley Vassilvitskii '21][Chen Silwal Vakilian Zhang '22] Primal dual bipartite matching
- [Davies Vassilvitskii Wang '24] Max-flow via Push Relabel

## **Prediction strategies for warm start algorithms:**

- [Khodak Balcan Talwalkar Vassilvitskii '22] Compete with best fixed prediction in hindsight
- [Blum Srinivas '25] Take advantage of weaker structure in sequences

# Three ideas

- (1) **Repeated Computations:** Sequences of related instances of a problem can be solved faster than one at a time
- (2) **Dynamic Algorithms/Data Structures:** Dynamic problems are easier with information about future updates
- (3) **Randomized Algorithms:** Randomized algorithms and data structures can be hedged to take advantage of extra information by incorporating a prior

# Dynamic Algorithms


**Example:** 3-vertex connectivity (triconnectivity) in dynamic graphs

## Dynamic model:

- Graph on  $n$  vertices undergoes edge insertions and deletions (fully-dynamic)
- Maintain data structure to efficiently query whether pairs of vertices are 3-vertex connected
- Minimize **amortized** time per update/query

## Challenge:

- Instance only changes a little, but solution can change a lot!
- Design data structures to “reuse” work to solve subproblems



# Example: Triconnectivity

Best known fully dynamic algorithm for triconnectivity:

$O(n^{2/3})$  (worst-case) update time [Galil Italiano Sarnak '99]

**Baseline:** better than recomputing solution from scratch every day

**Good:** update time sublinear in graph size

**Gold standard:**  $\text{polylog}(n)$  update time (exponential improvement over baseline)

**Compare to “offline dynamic” setting:**

- Sequence of updates and queries are given in one batch
- Algorithm returns answers to all queries at once

Best known **offline dynamic** algorithm for triconnectivity:

$\text{polylog}(n)$  (amortized) update time [Peng Sandlund Sleator '17]

with slick divide-and-conquer algorithm!

**Knowledge of future updates lets us reuse computation more efficiently!**

Gap between fully-dynamic and offline dynamic exists for many problems

# Imperfect Information

## Fully Dynamic Model

No information about future updates

## Predicted-Updates Model

Imperfect information about future updates

## Offline Dynamic Model

Perfect information about future updates

Fully dynamic is **pessimistic**, but offline dynamic is too **optimistic**...

Can we instead use an **imperfect prediction** of future updates?

**Learning-Augmented approach** [Liu Srinivas '24]:

- Ask for a (potentially infeasible) prediction of the input sequence
- Perform fully dynamic updates on the true sequence that may or may not be close to the predicted sequence

**Goal:**


- **Consistency:** when predictions are good, recover offline dynamic performance
- **Robustness:** always do at least as well as fully dynamic performance
- **Graceful degradation:** smooth tradeoff in error of prediction



# Under the Hood of Offline Dynamic Algorithms

## Divide-and-conquer approach:


- Associate each edge with “windows” for which it is **permanent**
- Each window is associated with at most twice as many edges as its size
- Process edges in batches from root to leaf, using **sparsifier** for problem



# What if the prediction is wrong?

Predicted  $e_1$  insertion on day 1, but happens on day 2 instead.

## Observation: errors are localized!




# What if the prediction is wrong? (Part 2)

Predicted  $e_4$  insertion on day 10, but happens on day 8 instead.

**Observation:** errors are localized!

...but not very localized, small errors can trigger large recomputations



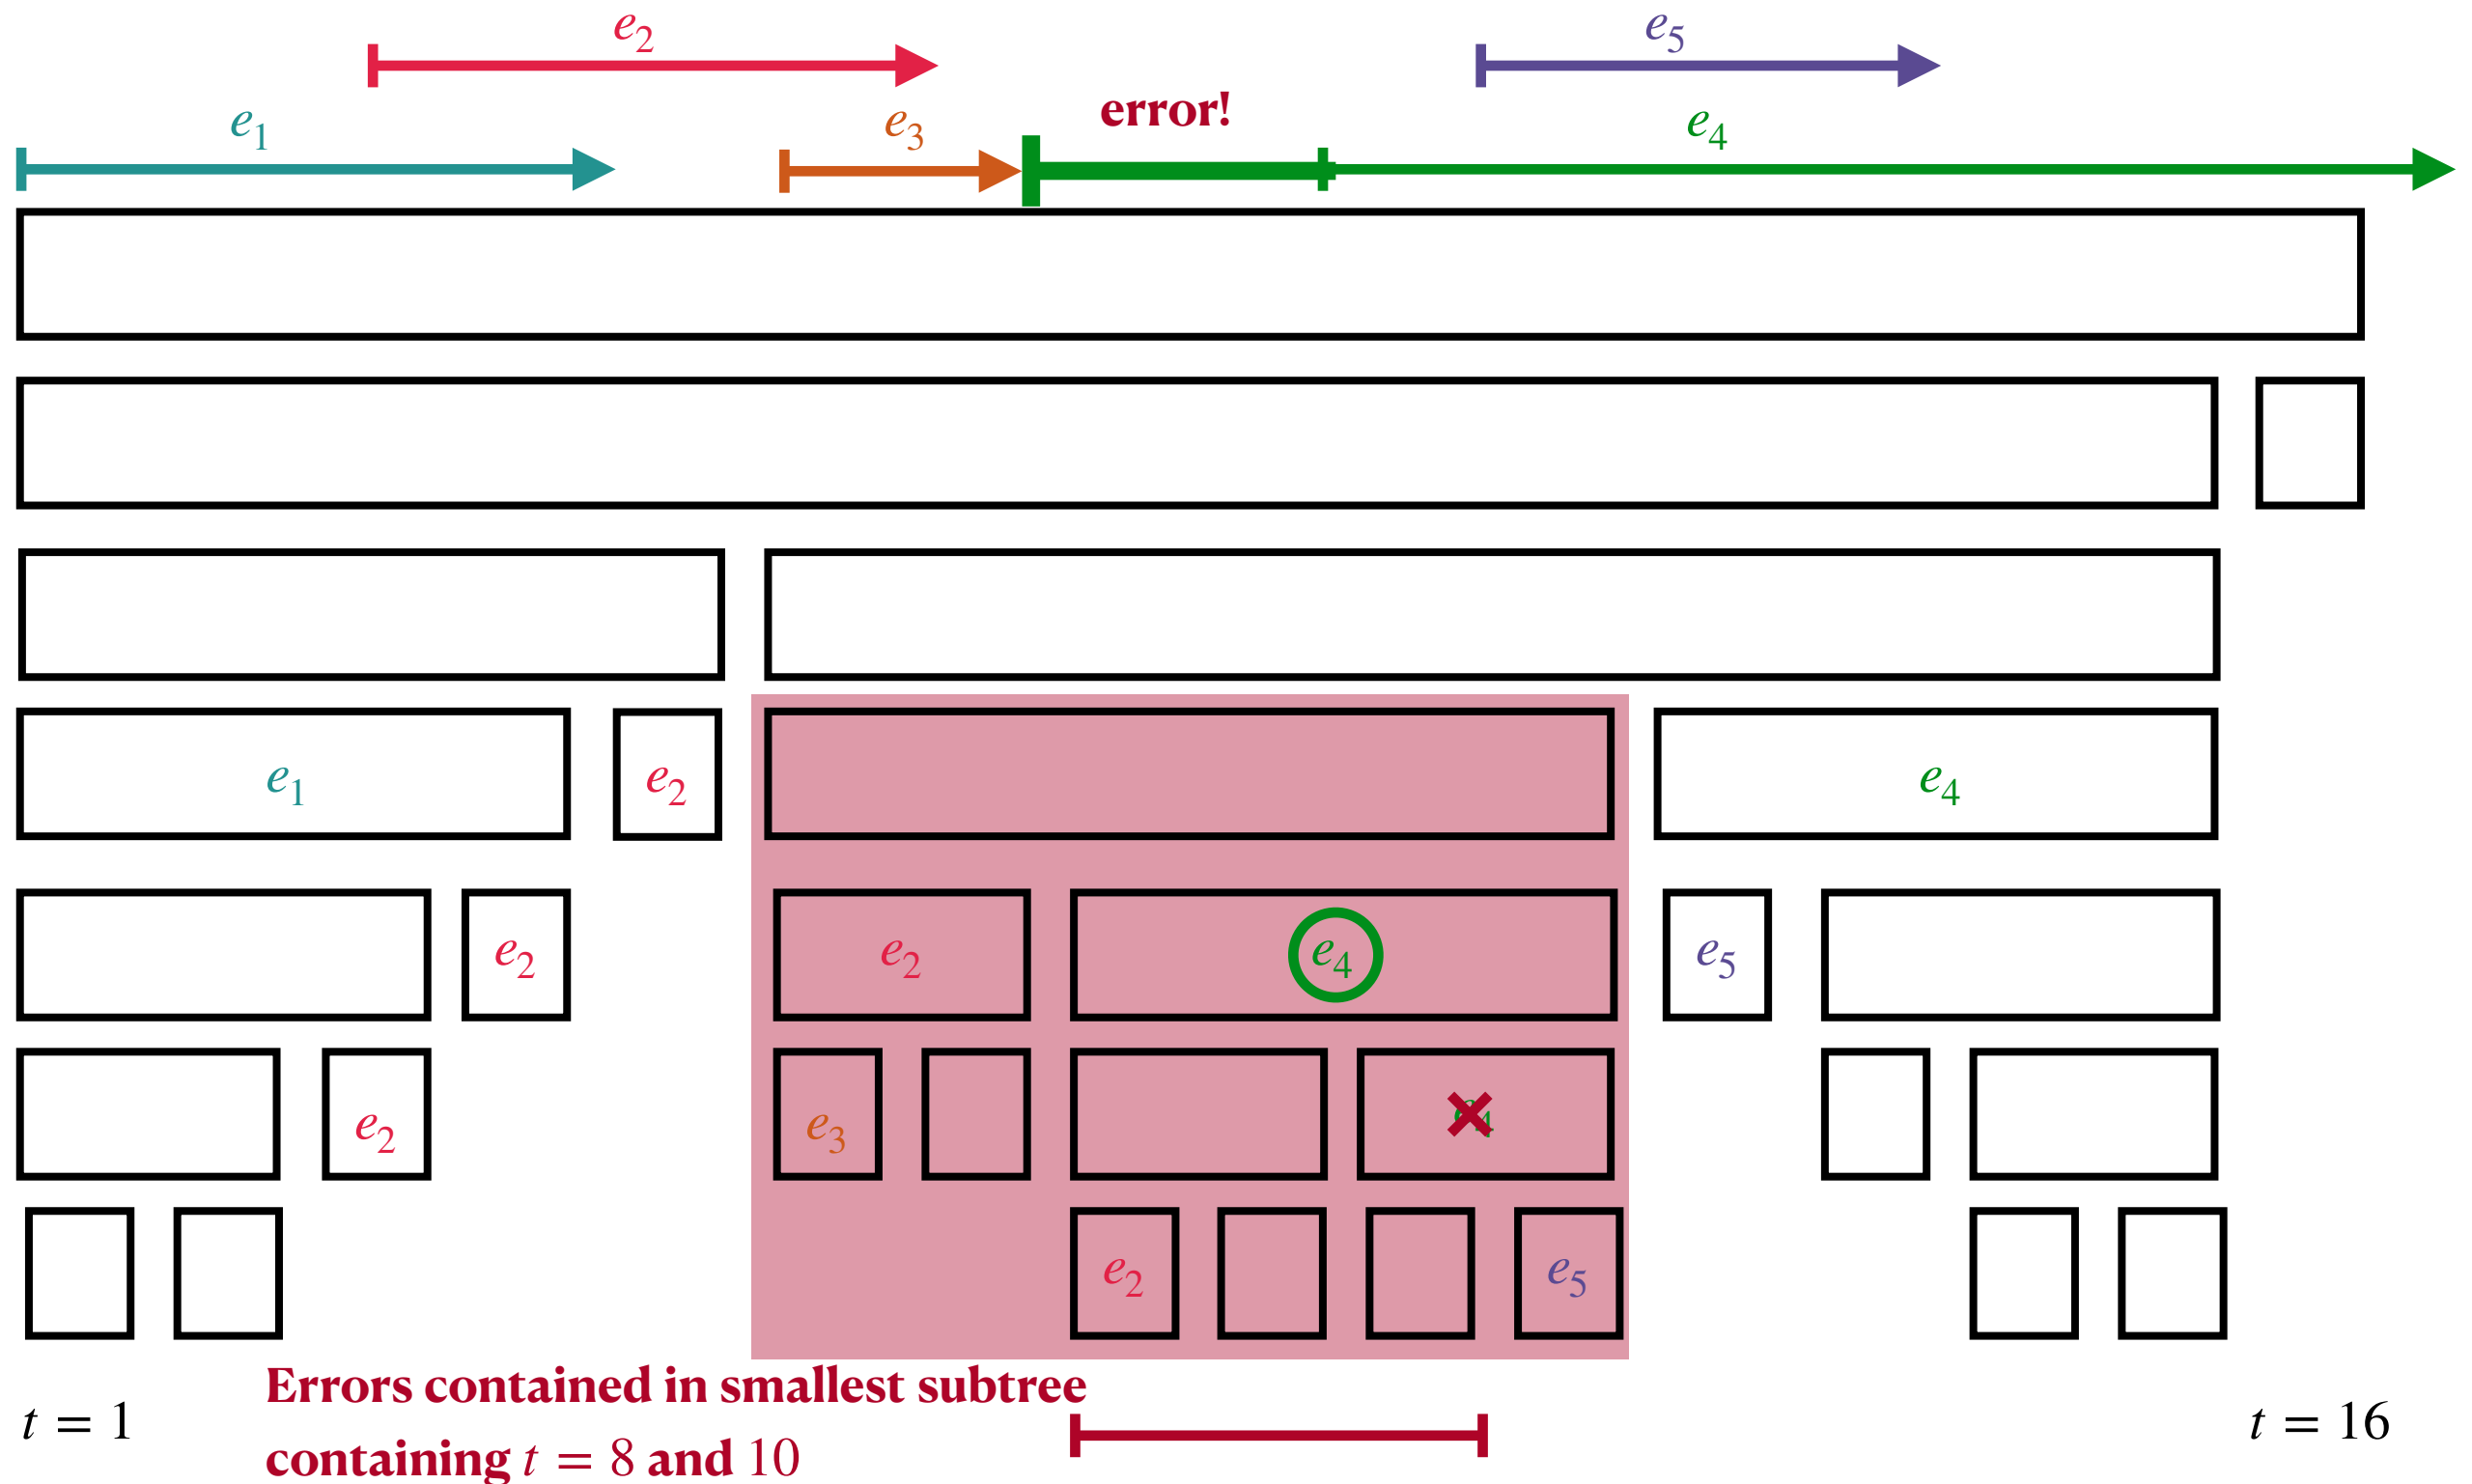
$t = 1$  Errors contained in smallest subtree  
containing  $t = 8$  and  $10$

$t = 16$

# Workaround

Use **random** divide-and-conquer tree

⇒ size of error subtree proportional to size of error



# Guarantees

Use this (and other) ideas to get an  
**fully-dynamic to offline reduction with predictions**

## **Informal Theorem** [Liu Srinivas '24]

Given a predicted sequence of events in advance, we can solve dynamic triconnectivity\* in total time over  $T$  updates

$$\widetilde{O}(\min\{T + \eta, T \cdot n^{2/3}\}),$$

Where  $\eta$  is the  $\ell_1$  error of the prediction (sum over events of the absolute prediction error in time).

\*Can do reduction for many problems with similar offline dynamic algorithms!

**Takeaway:** asymptotically lose **nothing** by taking advantage of predictions!

# Big Picture

Many recent works use similar and different techniques to use predictions to get improved dynamic algorithms:

- [van den Brand Forster Nazari Polak '23] Other graph and matrix problems
- [Agarwal Balkanski '24] Dynamic submodular maximization
- [McCauley Moseley Niaparast Singh '24] Incremental Topological Ordering
- [McCauley Moseley Niaparast Niaparast Singh '25] Approximate SSSP

Spiritually similar but **incomparable** to warm starts

## Dynamic Algorithms

- Structured input sequence (edge insertion/deletions)
- Small changes in input, potentially large changes in solution
- Can achieve sublinear update time

## Warm Starts

- Unstructured input sequence
- Potentially large changes in input, take advantage of similarities in solutions
- At bare minimum, must read input on each day  $\Rightarrow$  no sublinear update time

# References

**(2) Dynamic Algorithms/Data Structures:** Dynamic problems are easier with information about future updates

## Dynamic algorithms with predictions:

- [\[Liu Srinivas '24\]](#) Offline to online transformations
- [\[van den Brand Forster Nazari Polak '23\]](#) Graph and matrix problems
- [\[Agarwal Balkanski '24\]](#) Submodular maximization
- [\[McCauley Moseley Niaparast Singh '24\]](#) Incremental Topological Ordering
- [\[McCauley Moseley Niaparast Niaparast Singh '25\]](#) Incremental Approximate Single Source Shortest Paths

## Related ideas:

- [\[Peng Sandlund Sleator '17\]](#) Designing offline-dynamic algorithms
- [\[Peng Rubinstein '22\]](#) Fully-dynamic to incremental reduction with known deletion error

# Three ideas

- (1) **Repeated Computations:** Sequences of related instances of a problem can be solved faster than one at a time
- (2) **Dynamic Algorithms/Data Structures:** Dynamic problems are easier with information about future updates
- (3) **Randomized Algorithms:** Randomized algorithms and data structures can be hedged to take advantage of extra information by incorporating a prior



# Example: Treaps

Self-balancing tree data structure

**Goal:** Support element insertions, lookups, and deletions in  $O(\log(n))$  time,  $n$  is # of elements in the tree at any given time


**Rank**  $\in [0,1]$   
unif. random!

**Value**  $\in \mathbb{R}$

**Treap** = **Tree** + **Heap** [Aragon Seidel '89]

**Min-heap** w.r.t. rank: parent always has lower rank than children

**Search tree** w.r.t. value: left subtree contains only elements of smaller value, right subtree contains only elements of larger value



- $\log(n)$  depth with high probability (quicksort type analysis)
- Coupling between ranks and random tree allows dynamic updates

# Room for improvement?

- $\log(n)$  update/query time optimal for a **worst-case** sequence
- For a given update sequence, could minimize lookup time by storing frequently accessed elements closer to the root
  - E.g., splay trees rebalance with accesses to do this dynamically
  - Splay trees have conjectured dynamic optimality (within a constant factor of best self-balancing tree in hindsight for every access sequence)
  - Dynamic optimality: long standing open problem
- What if you had **predictions** about the access sequence in advance? Can you get a **provable** optimality guarantee, without sacrificing worst-case update time?

# Incorporate the prior

**Idea:** incorporate frequency information into the random rank! [Lin Luo Woodruff '23]

For an element  $x$ , let  $p(x)$  be the **predicted** frequency with which  $x$  will be accessed

**Rank**  $\in [0,1]$

**Value**  $\in \mathbb{R}$

**Standard Treap:** Choose rank  $\sim U[0,1]$

**Learning-Augmented Treap:** rank( $x$ )  $\sim -\lfloor \log \log(p(x)) \rfloor + U[0,1]$

**Hedge frequently accessed items to the top!**

**Provable guarantees** [Chen Cao Stepin Chen '25]:

- Static optimality, when predictions are perfect
- Worst-case guarantees from Treap, with tradeoff
- (Meets provable benchmarks for splay trees, when frequencies are estimated on-the-fly)

# Other Applications

**Takeaway:** Hedging is a natural idea. Learning-Augmented algorithms gives us the language to reason about it!

**Algorithmic challenge:** How do you design the **right** distribution?

- Caching [Lykouris Vassilvitskii '18]
  - Use frequency information to inform randomized cache eviction decisions
- Ski-Rental [Kumar Purohit Svitkina '18] ...
  - Use information about the future to make decisions in the present for online algorithms
- Min-cut [Moseley Niaparast Singh '25]
  - Use predictions about the min-cut to inform what vertices to contract in Karger-Stein

# References

**(3) Randomized Algorithms:** Randomized algorithms and data structures can be hedged to take advantage of extra information by incorporating a prior

## Learning-Augmented Search Trees:

- [Lin Luo Woodruff '23] Learning-Augmented Treaps with guarantees for stochastic accesses
- [Chen Cao Stepin Chen '25] Guarantees for general access sequences

## An incomplete list of other places to see this idea:

- [Lykouris Vassilvitskii '18] Caching
- [Kumar Purohit Svitkina '18], **and follow up work**, Ski-rental (proof of concept for online algorithms in general)
- [Moseley Niaparast Singh '25] Min-cut via Karger-Stein

# Conclusion

- (1) **Repeated Computations:** Sequences of related instances of a problem can be solved faster than one at a time
- (2) **Dynamic Algorithms/Data Structures:** Dynamic problems are easier with information about future updates
- (3) **Randomized Data Structures:** Randomized data structures can be hedged to take advantage of extra information by incorporating a prior

...and so much more!

**Takeaway:** Learning-Augmented algorithms gives us tools and frameworks to reason about interesting and practical new problems

Exciting time to join the field!

- Lots of great work over the last 5-10 years laying the groundwork
- Seeing the payoff in new results that take advantage of a new ways of algorithmic thinking
- Join us!

**THANKS!**

[vaidehi@u.northwestern.edu](mailto:vaidehi@u.northwestern.edu)