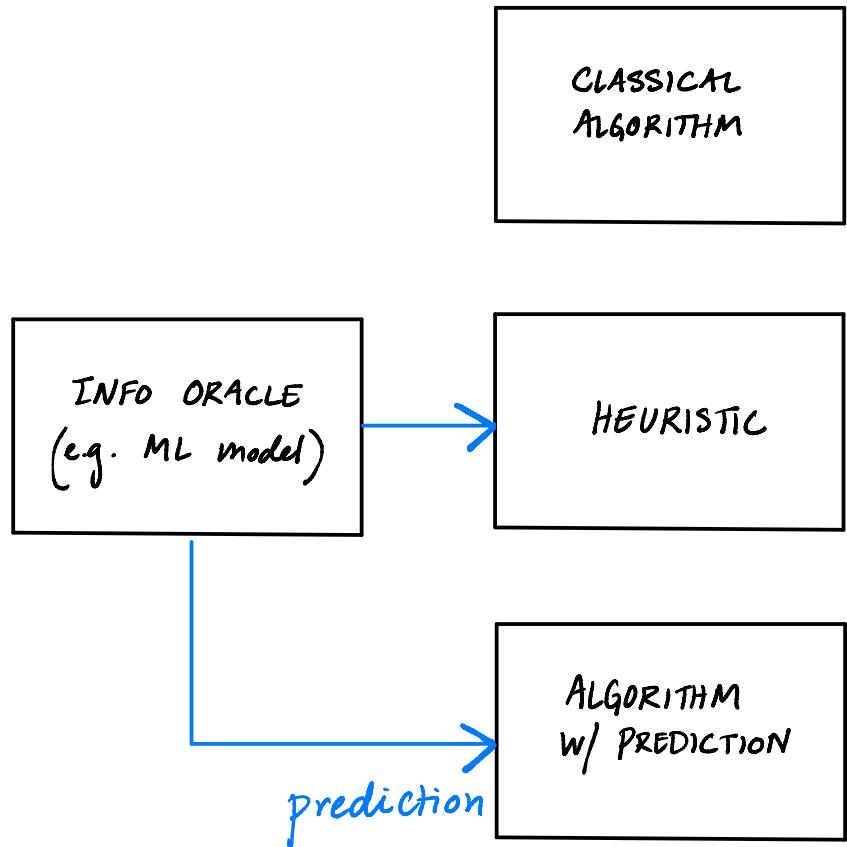


Competitive strategies to use "warm start" algorithms with predictions

Vaidehi Srinivas
Northwestern University

joint work with
Avrim Blum
Toyota Technological Institute, Chicago

ALGORITHMS WITH PREDICTIONS



Pessimistic:

Worst-case guarantee on every input

Optimistic:

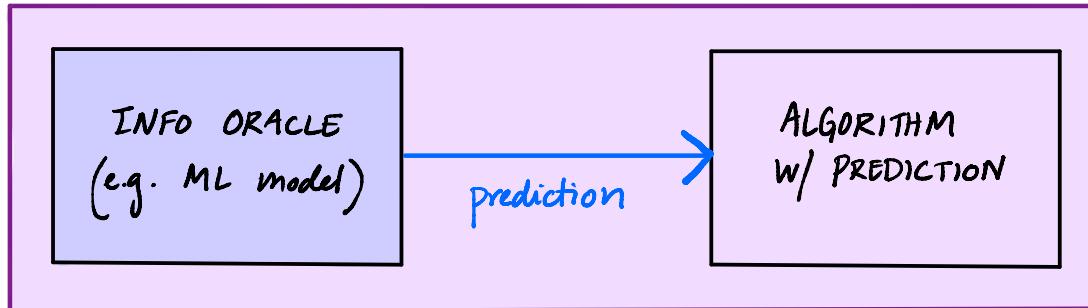
Performs well when information is good

Best of both worlds:

Good performance when info is good
and worst-case guarantee

LEARNING PREDICTIONS

Algorithm for sequences of instances



- ▶ Growing line of work in designing algorithms that use predictions

OUR GOAL: learn and use good predictions for algorithms with predictions

- ↳ "End to end" guarantees for solving sequences of related instances

WARM STARTS

Assume an instance space \mathcal{I} , metric solution space \mathcal{S}

Definition.

A warm start algorithm $A(I, P)$ for $I \in \mathcal{I}$, $P \in \mathcal{S}$
finds S^* (true solution for instance I)
in time at most $d(P, S^*)$.

Examples.

- ▶ Bipartite matching via learned duals

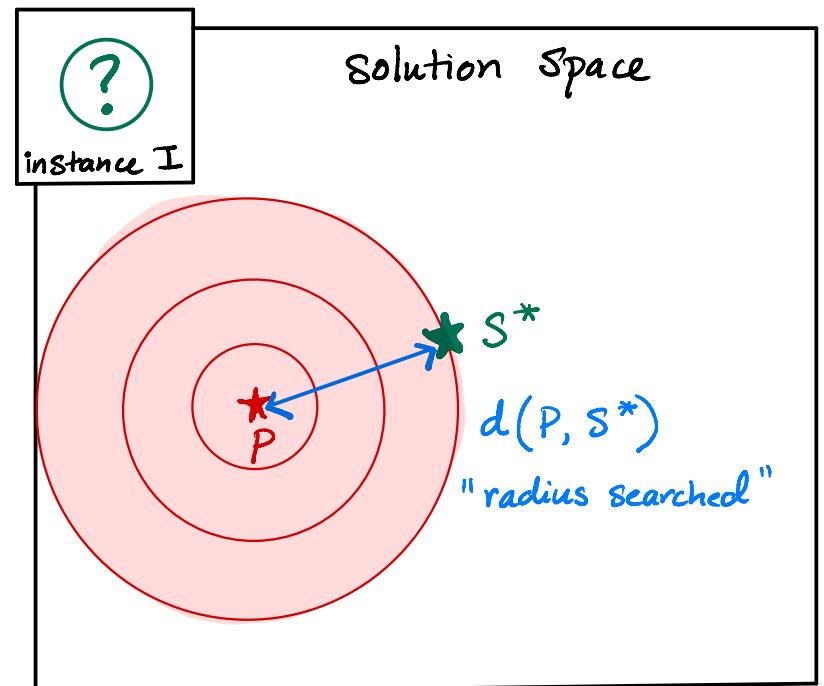
[Dinitz Im Lavastida Moseley Vassilvitskii 21],

[Chen Silwal Vakilian Zhang 22]

- ▶ Max Flow via Ford Fulkerson

[Polak Zub 22],

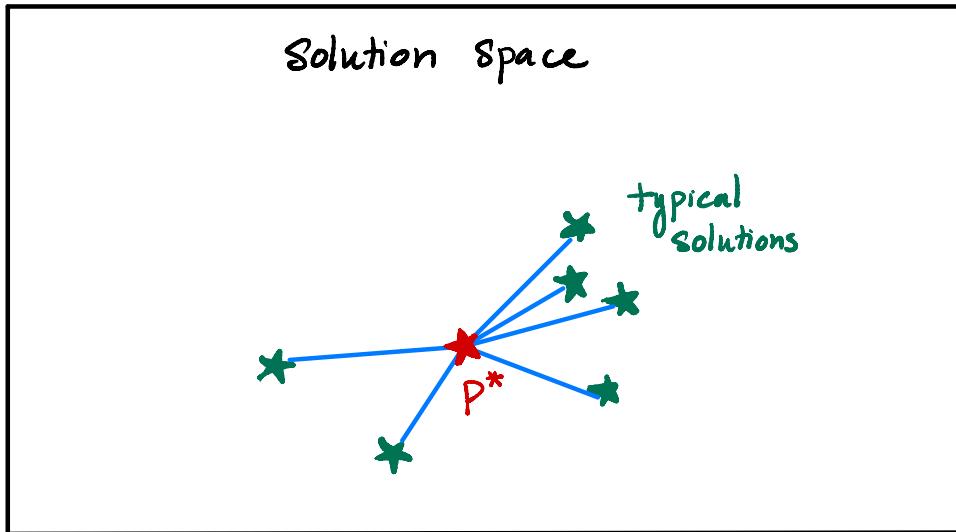
[Davies Moseley Vassilvitskii Wang 23]



HOW TO FIND PREDICTIONS: DISTRIBUTIONAL

Model: See instance-solution pairs drawn from a distribution

$$(I, s) \sim D \quad \text{over } I \times S$$



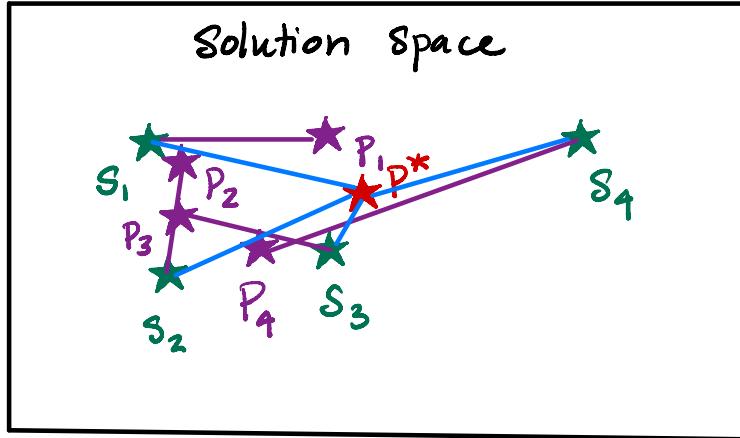
Strategy: Learn fixed prediction P^* to minimize

$$\mathbb{E}_{(I,s) \sim D} [d(P^*, s)]. \quad (\text{clustering cost})$$

On future instances, run $A(I, P^*)$.

HOW TO FIND PREDICTIONS: ONLINE

Model: A sequence of instances I_1, \dots, I_t arrive online



Strategy: Select predictions to compete with best fixed prediction
in hindsight p^* via online convex optimization
[Khodak Balcan Talwalkar Vassilvitskii '22]



Distributional and online guarantees require solutions to be very close together to provide strong guarantees



Can we get strong guarantees in more general settings?
↳ still require some structure!

WARM UP: COMPETING WITH k PREDICTIONS (DISTRIBUTIONAL)

? What if there are multiple types of instances?

Strategy: Learn predictions P_1, \dots, P_k to minimize expected distance to some prediction

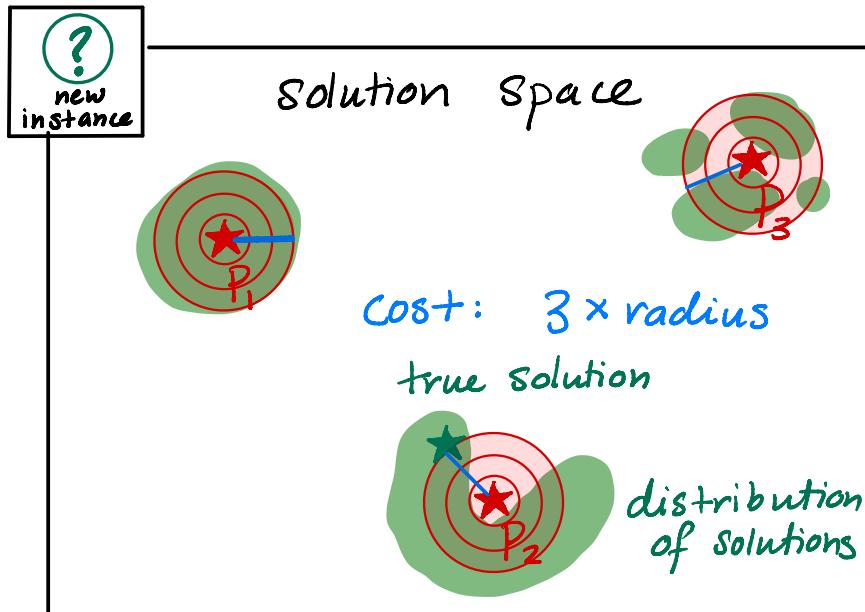
$$\text{minimize}_{(I, S) \sim D} \mathbb{E} \left[\min_i d(P_i, S) \right] \quad (\text{k-medians clustering cost})$$

Algorithm.

For a fresh $(I, S) \sim D$,

run $A(I, P_1), \dots, A(I, P_k)$
"in parallel"

output computation that completes first.



Guarantee.

Expected work of algorithm is $O(k)$ times the k -median clustering cost.

OVERVIEW OF RESULTS

? Can we get around the $O(k)$ factor by using other structure?

- ▶ Show how to learn a mapping from instances to predictions, $h: \mathcal{I} \rightarrow \mathcal{S}$
- ▶ Get $O(1)$ approximation guarantee when h comes from a learnable class \mathcal{H}

? What can we do in the online setting?

- ▶ Formulate the "online ball search" problem
- ▶ Provide an algorithm that works well when solutions are k -clusterable in hindsight.

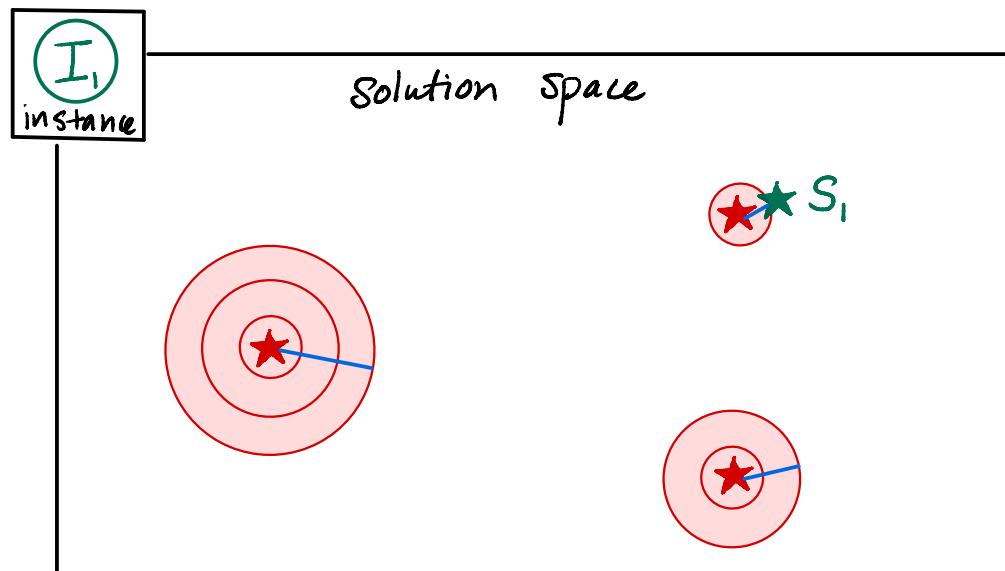


ONLINE FORMULATION

Problem.

- ▶ Have access to warm-start algorithm A (search function)
- ▶ On each day $t \in [T]$, instance I_t arrives online
 - ▶ Algorithm chooses predictions and search schedule to find solution S_t of I_t

Goal: minimize total radius searched over all days



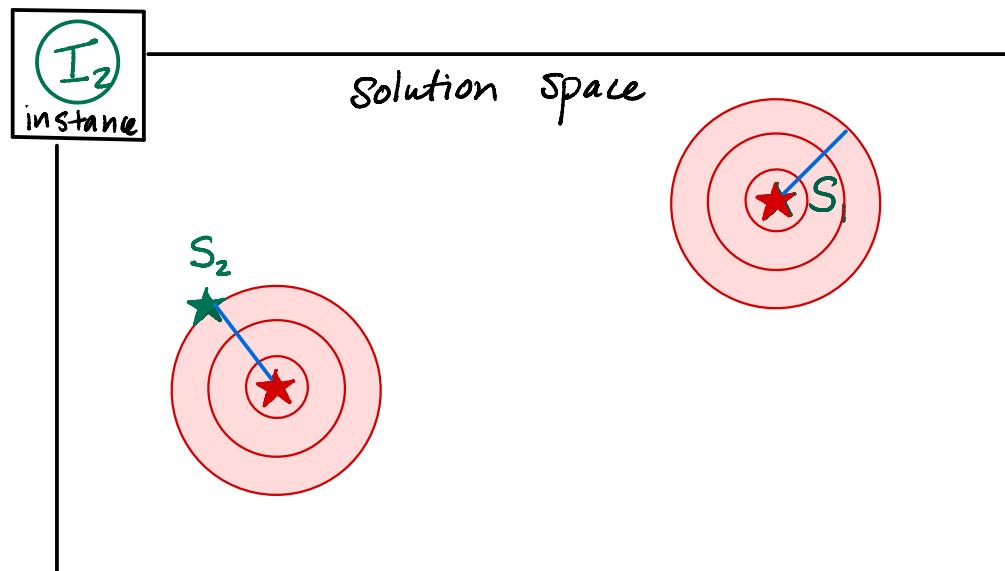
total radius searched:
6

ONLINE FORMULATION

Problem.

- ▶ Have access to warm-start algorithm A (search function)
- ▶ On each day $t \in [T]$, instance I_t arrives online
 - ▶ Algorithm chooses predictions and search schedule to find solution S_t of I_t

Goal: minimize total radius searched over all days



total radius searched:

$$6 + 6 = 12$$

TYPICAL BASELINES

WANT: Analyze algorithms via competitive analysis

- ⚠ Offline optimum guesses solution perfectly every day and incurs 0 cost!
- ▶ Analysis does not assume enough structure in the sequence

Previous approach: [Khodak Balcan Talwalkar Vassilvitskii '22]

Competes with best fixed prediction in hindsight

↳ restricts power of the offline strategies that we compete against

Modeling contribution: What is a more powerful set of offline strategies that we can design competitive algorithms against?

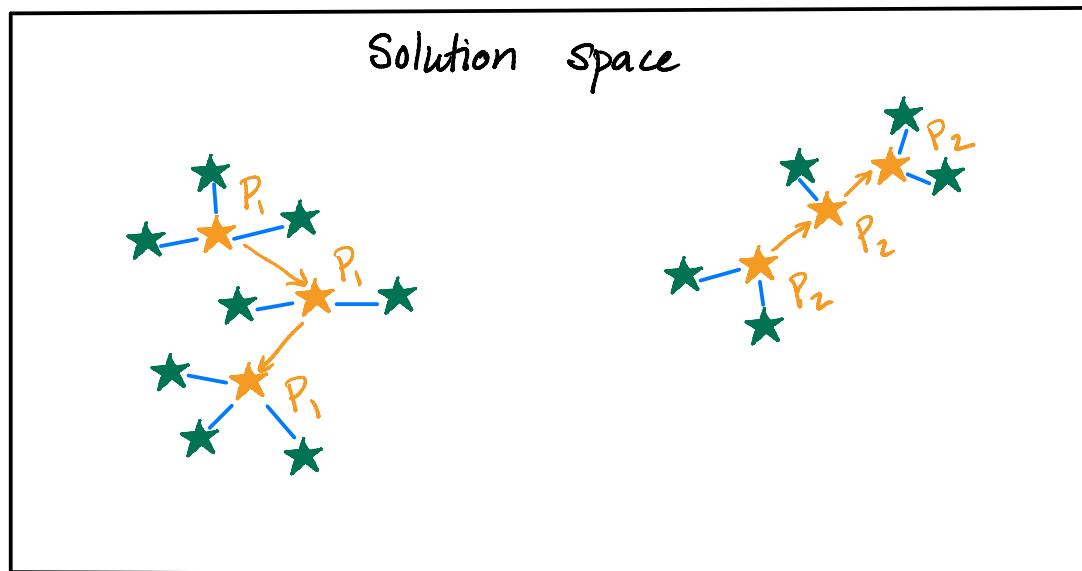
BASELINES

Definition [k offline trajectories]

Offline strategy that maintains a set of k predictions, and on each day pays

hit cost: distance from today's solution to closest prediction

movement cost: distance moved by all predictions



Similar to offline
k-server cost

- ▶ Sequences that are "clusterable" have good baselines
- ▶ WANT: Design algorithm that is competitive with baselines

ONLINE RESULT

Theorem [Blum S. 25].

We give an algorithm that is $O(k^4 \ln^2 k)$ - competitive with any set of k offline trajectories in radius searched.

Furthermore,

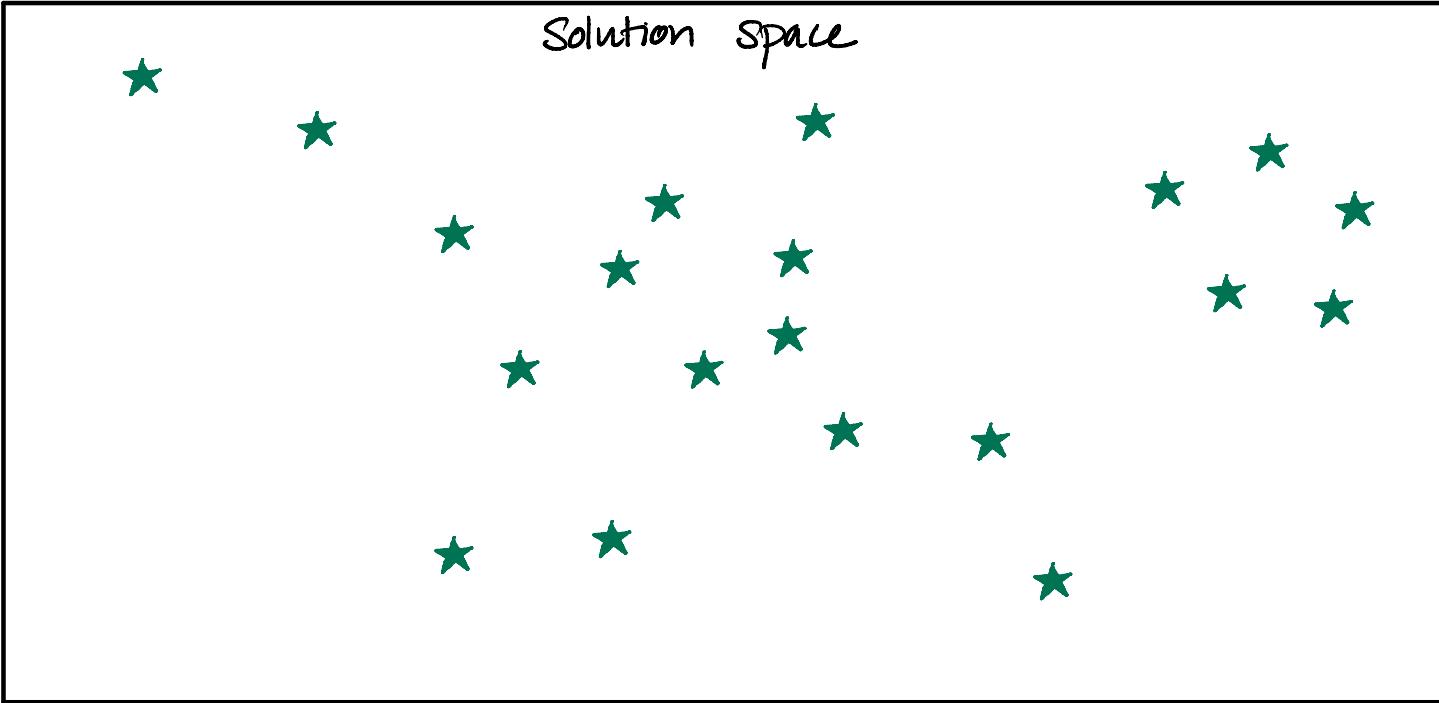
- ▶ total runtime of algorithm is bounded by $O(1)$ times the radius searched

- ▶ algorithm is deterministic (robust to adaptive adversary)

- ▶ guarantee holds for all k simultaneously

- ▶ Our result + warm start algorithm \Rightarrow end-to-end guarantee for sequences of instances that display "clustering" structure

ONLINE CHALLENGES

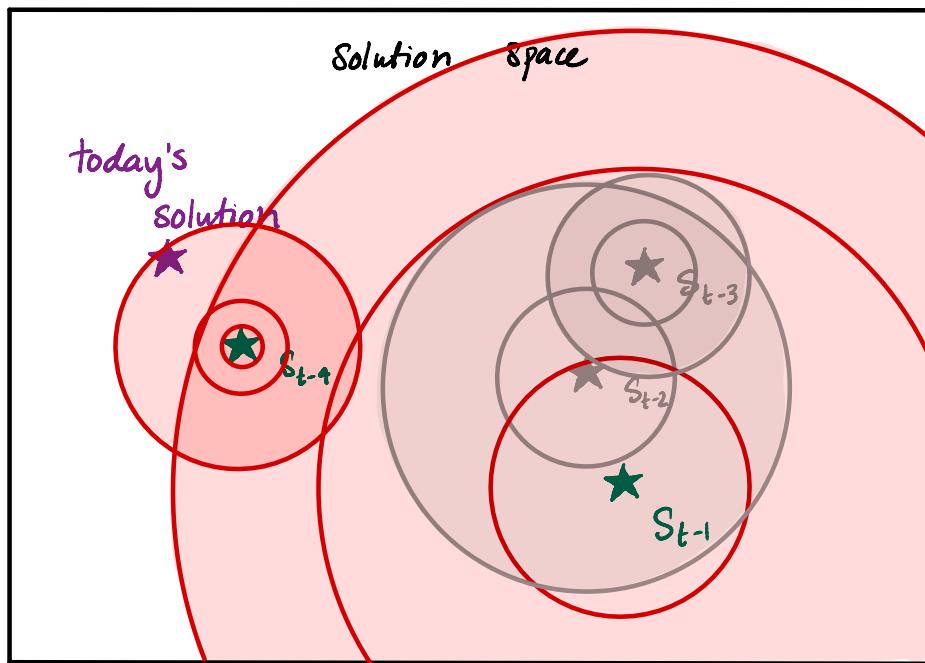


- ▶ Competing with 1 trajectory is simple!
Use yesterday's solution as today's prediction

Issues with multiple trajectories:

- ① Don't know which trajectory today belongs to
- ② Don't know how to group previous solutions into trajectories

OVERVIEW OF ONLINE ALGORITHM



Algorithm. [Quadratic decay]

On day t :

- ▶ Search from i^{th} most recent solution at rate $\frac{1}{i^2 \ln^2 i}$
- ▶ if Solution i "subsumes" Solution j :
kill solution j and increase rates
of slower solutions

- ▶ WANT: Most recent Solution from same trajectory will find *
- ▶ If the algorithm runs long enough, other trajectories "collapse" and solution from same trajectory elevated to high rate
- ▶ Spiritually related to double coverage algorithm for k -server

CONCLUSION

We study how to learn and use predictions to solve sequences of related instances using a "warm start" algorithm.

Distributional setting:

- ▶ a simple strategy pays $O(k)$ times as much as the k -clustering cost of the distribution
- ▶ Can potentially avoid the $O(k)$ -factor blow-up by learning a mapping from instances to solutions

Online setting:

- ▶ Can design an algorithm to solve sequences of instances that are k -clustered in hindsight.

THANKS!

vaidehi@u.northwestern.edu