

# THE PREDICTED-UPDATES DYNAMIC MODEL: OFFLINE TO FULLY-DYNAMIC TRANSFORMATIONS

*Dynamic Algorithms with Predictions*



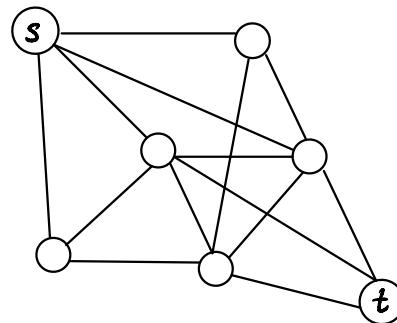
Quanquan C. Liu  
Yale



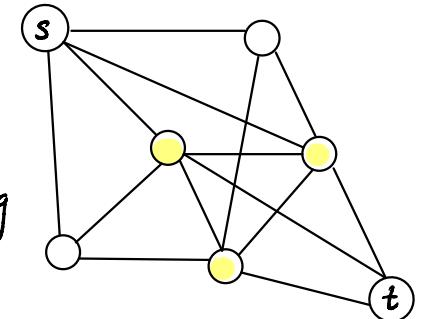
Vaidehi Srinivas  
Northwestern University

## CASE STUDY : TRICONNECTIVITY

updates  
edge insertions and  
deletions



queries  
"can I disconnect  
s and t by deleting  
3 vertices from the  
graph?"



Best known fully-dynamic algorithm:  $\tilde{O}(n^{2/3})$  update time

[Galil Italiano Sarnak '99]

? Is the barrier information? or computation?

A Information! Can solve offline dynamic problem

in  $\tilde{O}(\frac{T}{\sqrt{t}} \text{ polylog } n) = \tilde{O}(T)$  time. [Peng Sandlund Sleator '17,  
total # of updates "  $\tilde{O}(1)$  per update" Holm Rotenberg '20]

- Solve with slick divide-and-conquer algorithm
- Candidate for "algorithms with predictions"

# ALGORITHMS WITH PREDICTIONS

"prediction"

Goal: effectively utilize extra untrusted information about an instance  
e.g. "today's instance is likely similar to yesterday's instance"

Modeling challenge: What prediction should we ask for?

How should we measure prediction error?

Is it reasonable to get this in practice?

Ideal guarantee: [Lykouris Vassilvitskii '21]  not always possible!

- ① Consistency: if the prediction is good, do better than the worst case
- ② Robustness: if the prediction is bad, achieve a worst-case guarantee  
a.k.a. "competitiveness"
- ③ Graceful degradation: smooth tradeoff between ① and ②  
"Robustness", "smoothness"

Examples: Ski-rental

[Kumar Purohit Srivastava '18]

pred: # of days able to ski

Max-Flow via Ford-Fulkerson  
[Davies Moseley Vassilvitskii Wang '23]

pred.: solution to max-flow instance  
"warm start"

Example:

## THE (STANDARD) DYNAMIC MODEL

Goal: compute  $f(G)$ , for dynamic graph  $G$

Initially : start with empty graph

On each day  $t \leq T$ : "time horizon"

one edge inserted or deleted from  $G$

algorithm outputs  $f(G_t)$   $\xrightarrow{T}$  state of graph on  $G_t$



algorithm must always output correct solutions!



objective: minimize update time  $\xrightarrow{\text{runtime per update}}$

either      worst-case per update time

or            amortized update time (average over all updates)

$G_0$ : . . .

$G_1$ : + . .

$G_2$ : . + .

$G_3$ : - . .

# NEW! PREDICTED - UPDATES DYNAMIC MODEL

Goal: compute  $f(G)$ , for dynamic graph  $G$

Initially : start with empty graph,  
receive set of  $T$  predicted updates :

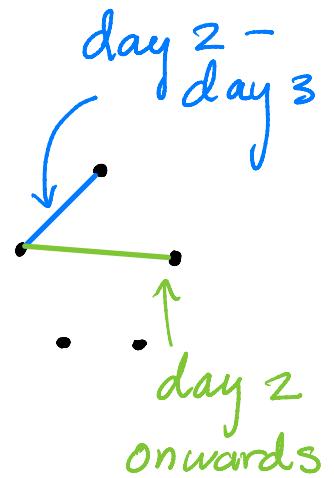
(update, predicted day)  $\leftarrow$  no feasibility requirement!

On each day  $t \leq T$ :

one edge inserted or deleted from  $G$

algorithm outputs  $f(G_t)$  true state of graph on  $G_t$

Example  
prediction:



⚠ algorithm must always output correct solutions!

⚠ runtime of algorithm may depend on  
pred. error magnitude

$$\ell_1\text{-error}(\text{predictions}) = \sum_{\text{updates}} \left| \frac{\text{predicted update time}}{\text{true update time}} - 1 \right|.$$

# OFFLINE TO FULLY-DYNAMIC TRANSFORMATION

Theorem [Liu S. '23] (Informal).

Given an offline divide-and-conquer\* algorithm to compute  $f(\cdot)$  that does  $\tilde{O}(T)$  work,

↗ "backstop"

and a fully-dynamic algorithm,  $\underline{B}$ , with update time  $\text{update}_B$ , we can design a predicted-updates dynamic algorithm that does total work

$$\tilde{O}\left(\min\left\{\frac{T}{\text{update}_B}, \frac{\|\text{prediction error}\|_1}{\text{update}_B}\right\} + T \cdot \text{update}_B\right)$$



Consistency      graceful degradation      robustness

\* we require the divide-and-conquer to be "balanced," and satisfy a technical condition

## INTERPRETATION

Work:

$$\tilde{O}\left(\min\left\{\underbrace{T}_{\text{consistency}} + \frac{\|\text{prediction error}\|_1}{\text{graceful degradation}}, \underbrace{T \cdot \text{update}_B}_{\text{robustness}}\right\}\right)$$

As long as average prediction error of updates  $\leq \text{update}_B$ , we beat the fully dynamic backstop.

<u>Applications</u> :	Triconnectivity	$\tilde{O}(n^{2/3}) \rightarrow \tilde{O}(1)$
	$k$ -Edge connectivity	$n^{o(1)} \rightarrow \tilde{O}(1)$
	Minimum Spanning Forest	$\tilde{O}(1) \rightarrow \tilde{O}(1)$ , but simpler!

Extends to incremental and decremental settings!

See also [van den Brand Forster Nazari Polak '24]

# ALGOS W/ PREDs. : A LAY OF THE LAND

## Warm starts (static)

- + consistency AND robustness simultaneously
- Predictions typically large  
(encode problem solutions)

## Online algorithms

- Inherent tradeoff between consistency and robustness
- + Predictions typically small  
(encode future updates)



## Dynamic algorithms

- + consistency AND robustness simultaneously
- + Predictions are small  
(encode future updates)

### Examples:

Ski-rental  
[Kumar Purohit Srivastava '18]

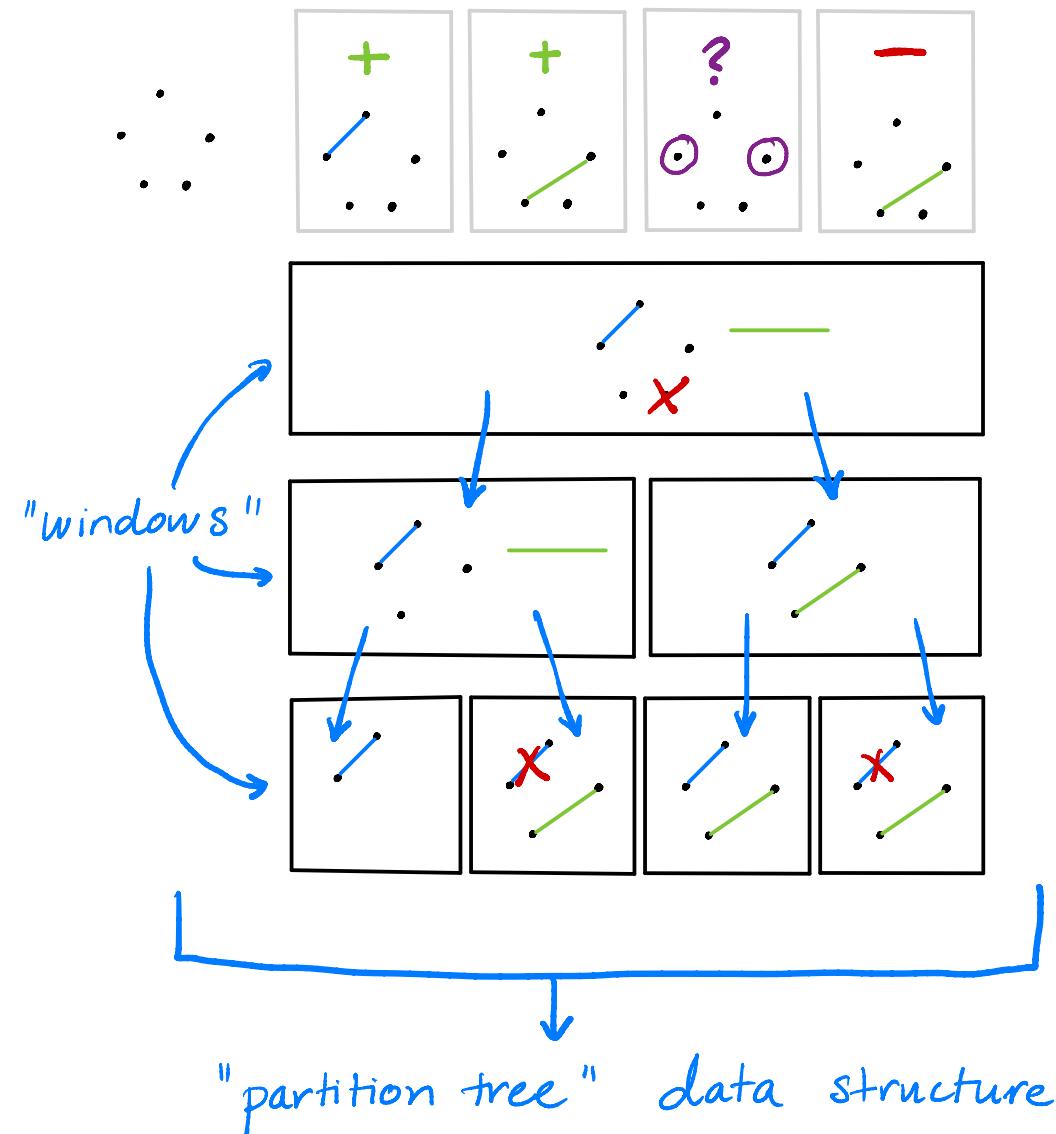
pred: # of days able to ski

Max-Flow via Ford-Fulkerson  
[Davies Moseley Vassilvitskii Wang '23]

pred.: Solution to max-flow instance  
"warm start"

# DIVIDE-AND-CONQUER OFFLINE DYNAMIC ALGORITHMS

Example : connectivity



## Principles:

### ① contraction:

Contract structures

that are "permanent"  
for a window

### ② reduction :

remove structures

that are irrelevant for  
a window

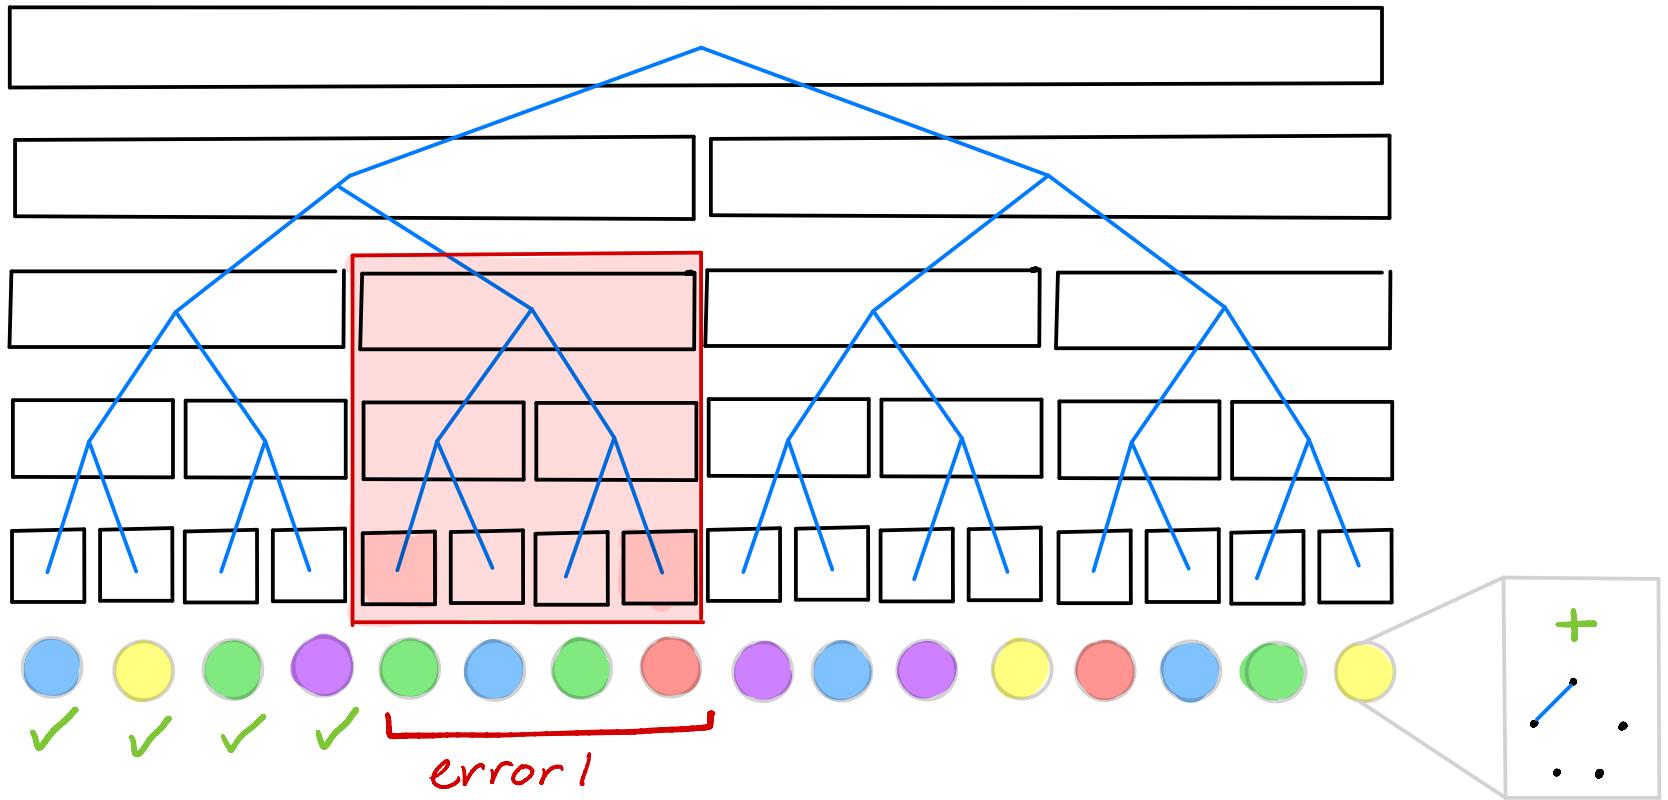
build a Sparsifier for your  
problem

## Goal:

work done per window is  
linear in the # of events

## A FIRST ATTEMPT

Run offline algorithm using predictions in  $O(T \log T)$  work.



Observe realized updates:



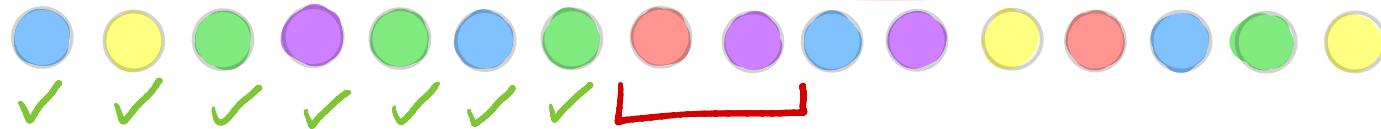
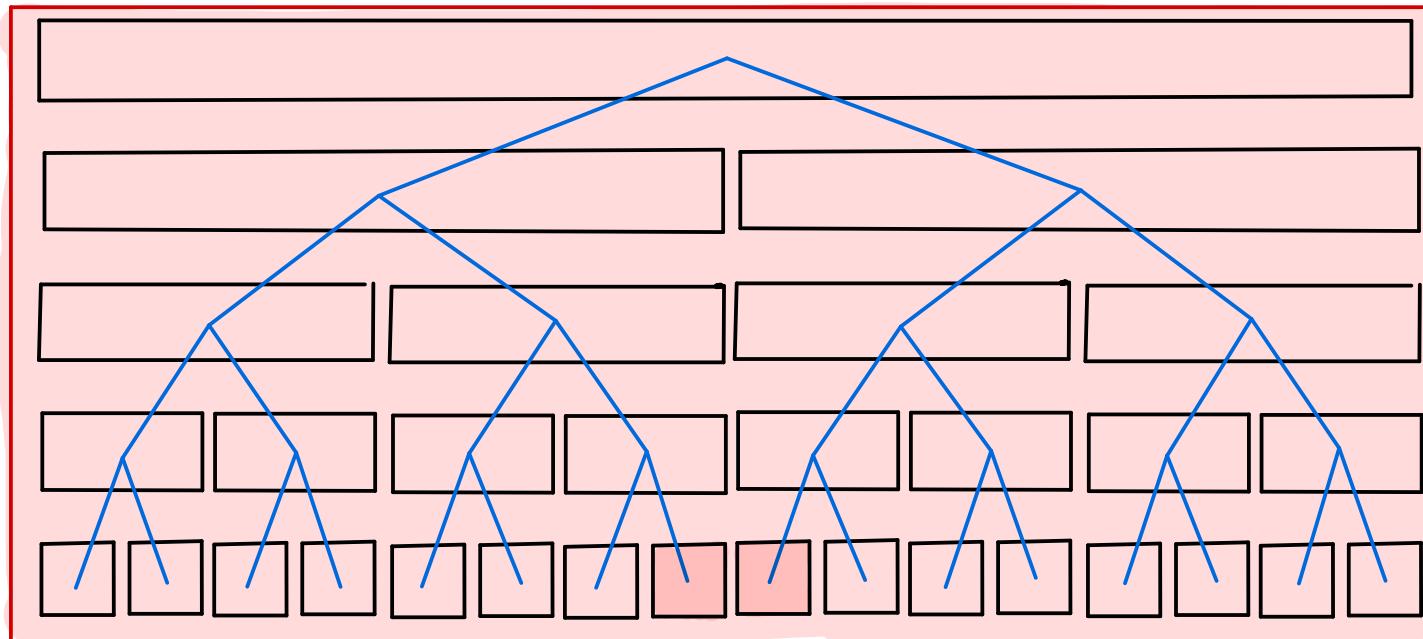
✗ Precomputed solutions incorrect

✓ Only need to fix small subtree!

Note: area of = work to recompute 9

## POTENTIAL PITFALL

Run offline algorithm using predictions in  $O(T \log T)$  work.



Observe realized updates:



error!

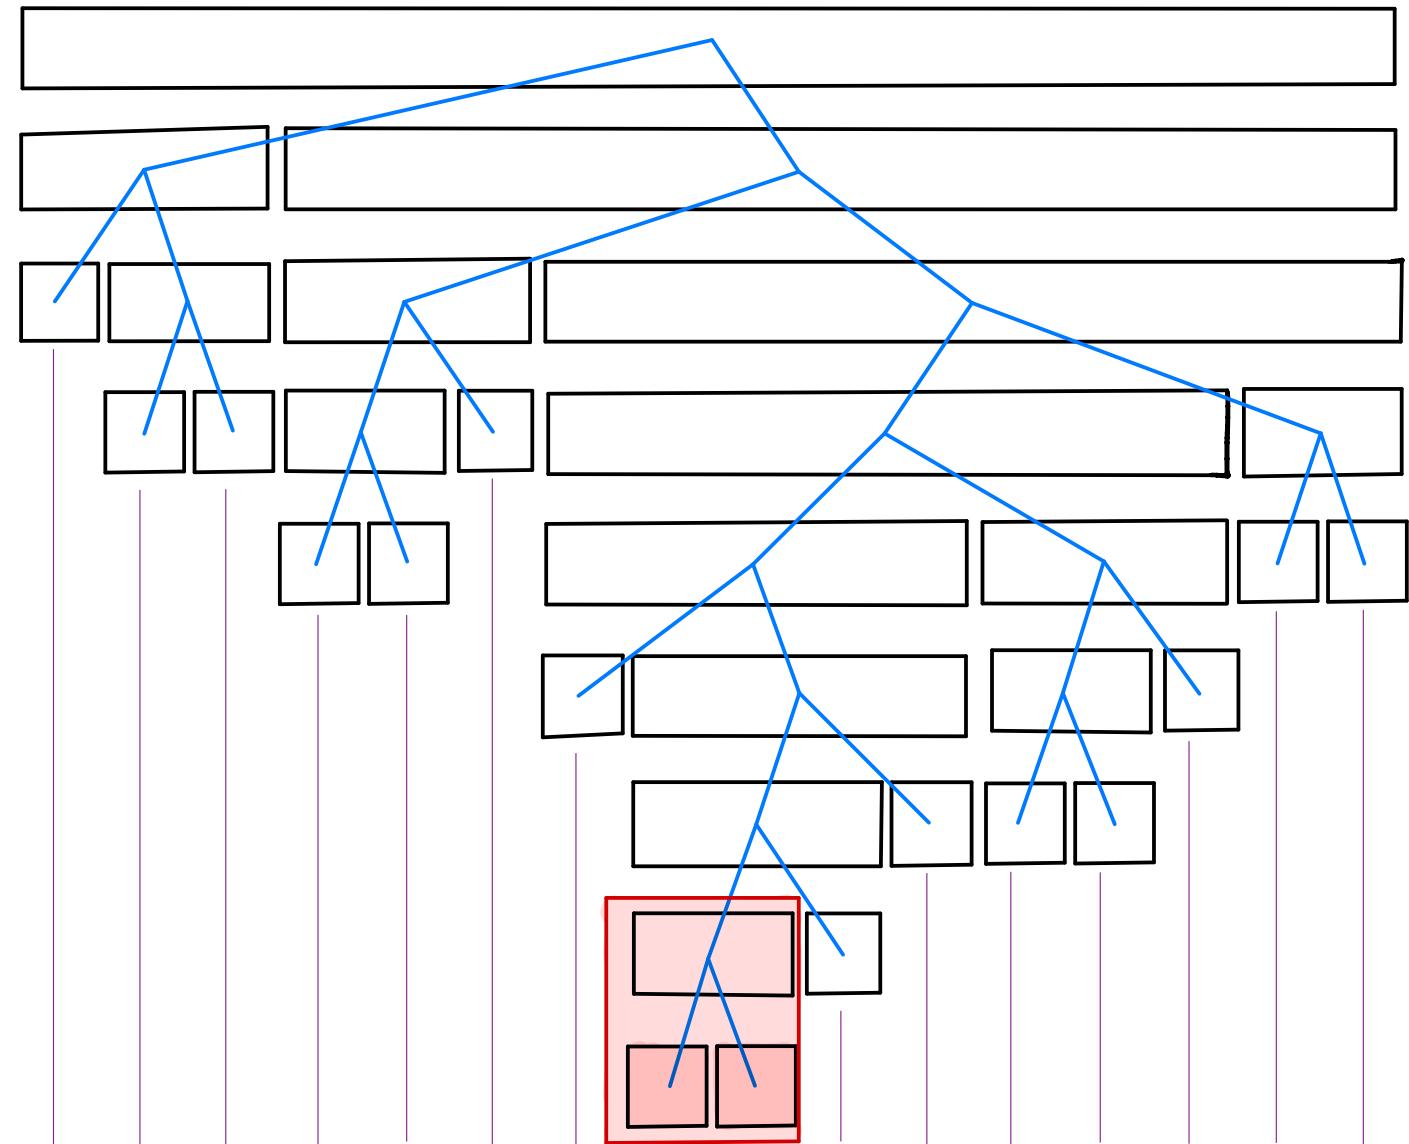
area of =  
work to recompute

➊ Small error triggers large recomputation!

# SOLUTION: RANDOM PARTITION-TREE



divide-and  
- conquer  
still efficient!



predicted:



realized:

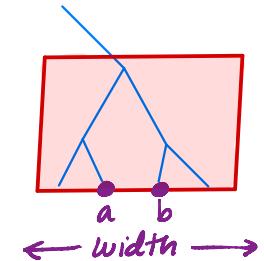


# PROPERTIES OF RANDOM PARTITION TREES

Property 1: Random trees preserve distance in expectation

$$\mathbb{E} \left[ \begin{array}{c} \text{Size of lowest common ancestor} \\ \text{of } a \text{ and } b \end{array} \right] = O(|b-a| \log T)$$

→ Randomized tree-embedding for the line metric  
see e.g. [Fakcharoenphol Rao Talwar '04]



Property 2: Random trees are low depth.

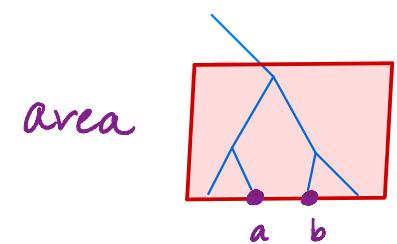
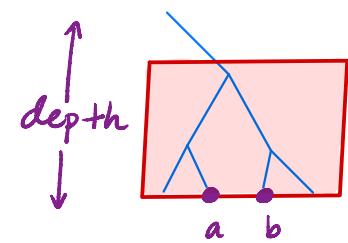
$$\mathbb{E} [\text{depth of random tree}] = O(\log T)$$

→ Divide-and-conquer is efficient

Implication: The expected work to recompute the smallest subtree containing  $a$  and  $b$  is

$$O(|b-a| \log^2 T).$$

→ work to fix error linear in magnitude of error!



## WHAT DID WE SWEEP UNDER THE RUG?

- ✓ Early update problem : (event occurs before predicted time)
  - ▶ handled by random partition tree

- ✓ Late update problem : (event occurs after prediction)
  - ▶ use doubling search, reduce to early update problem

⚠ each update can contribute  $O(\log T)$  events over lifetime

⚠ predictions can assign many events to the same day  
(predictions can be infeasible)

- ✓ Overscheduling problem : Divide-and-conquer breaks if too many events on one day!

- ▶ use online matching to make predictions feasible [Gupta Lewi '12]
- ▶ relax data structure to handle  $O(\log T)$  events per day

# HEY, WHAT ABOUT ROBUSTNESS ?

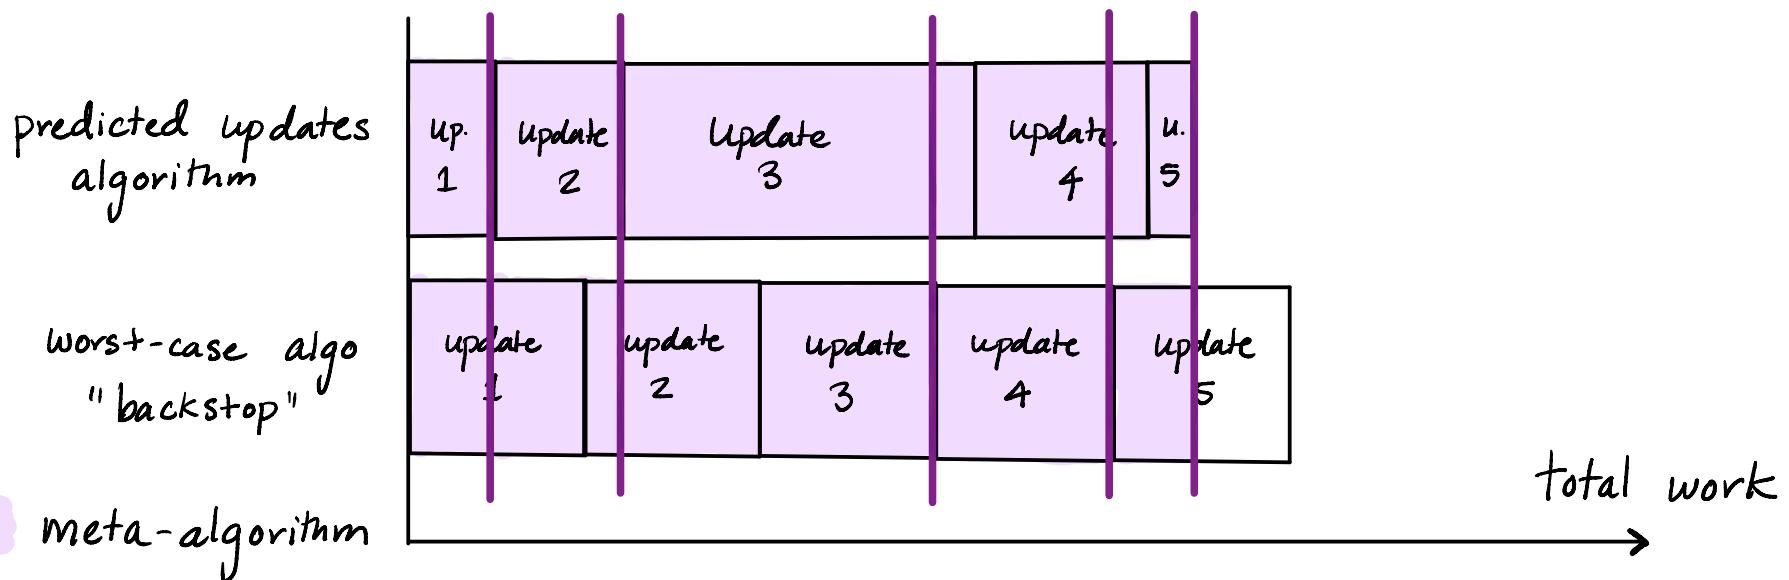
Work to handle  $T$  updates :

$$\tilde{O}\left(\min\left\{\frac{T + \|\text{prediction error}\|_1}{\text{update}_B}, \frac{T \cdot \text{update}_B}{\text{update}_B}\right\}\right)$$

saw how to do this

know how to do this

? How to get both simultaneously?



- ✓ also allows us to boost expectation bounds to high probability

## TAKEAWAYS

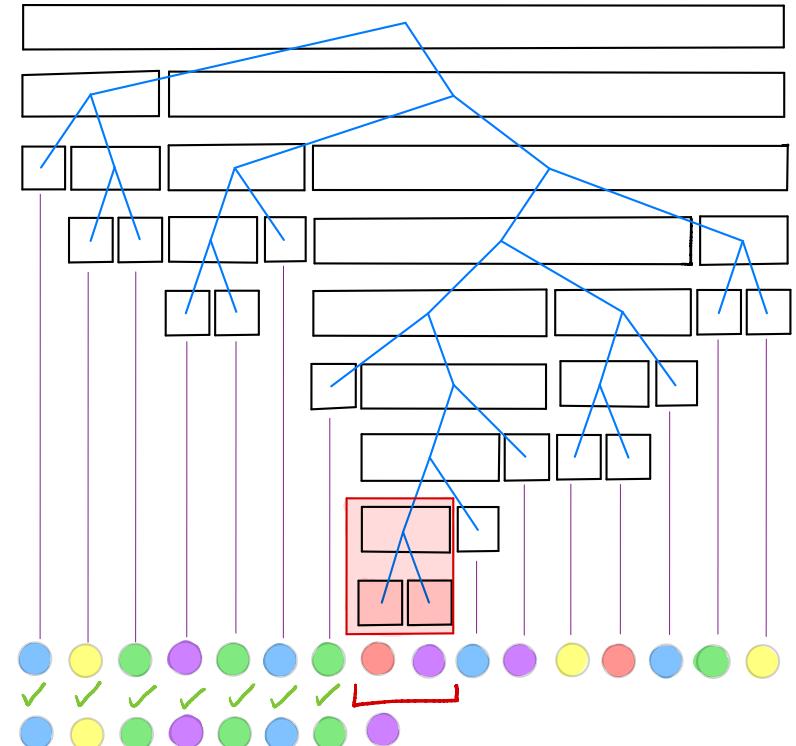
Work to handle T updates:

$$\tilde{O}(\min \{ T + \|\text{prediction error}\|_1, T \cdot \text{update}_B \})$$

- ① For many dynamic problems, can take advantage of extra information "for free!"
  - ▷ algorithms with predictions
- ② New model smoothly interpolates between offline-dynamic and fully-dynamic settings
  - ▷ "beyond-worst-case" models
- ③ Random partition-tree is simple and powerful
  - ▷ new theoretical insight into lifting offline and partially-dynamic algorithms to "fully-dynamic setting"

## FUTURE DIRECTIONS

- ▷ Better dependence on error?  
perhaps  $\ell_0$ ?
- ▷ Make deterministic or otherwise robust to adaptive adversaries
- ▷ Worst-case per update vs. amortized bounds
- ▷ Better dynamic subroutines for static problems



Dynamic algorithms with predictions is taking off!

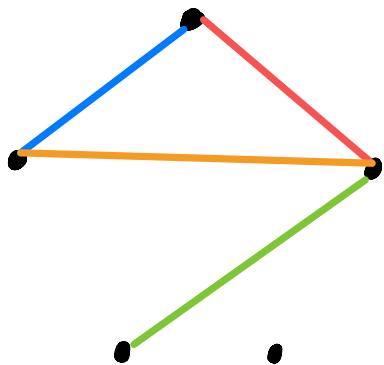
[van den Brand Forster Nazari Polak '24], [Henzinger Saha Seybold Ye '23],  
[Agarwal Balkanski '23]

THANKS!

Vaidehi@u.northwestern.edu

## EXTENSION TO PARTIALLY-DYNAMIC SETTINGS

Stronger assumption: incremental algorithm (vs. offline dynamic)



↳ dynamic algorithm that only handles  
edge insertions

worst-case per update time

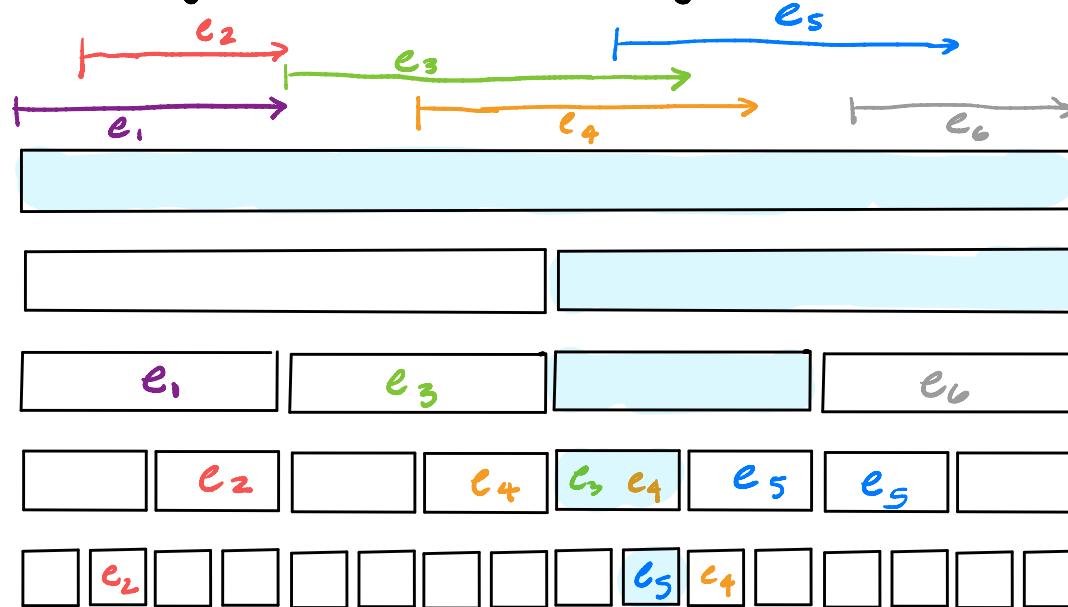
Stronger guarantee: only need predictions of deletion times

elements arrive arbitrarily, inserted  
element announces predicted deletion time

Same framework extends to this setting!

# INCREMENTAL TO FULLY-DYNAMIC TRANSFORMATION

First, design offline dynamic algorithm using incremental algorithm A



At each window, insert "permanent edges" of that window

Total work:  $\tilde{O}(T \cdot \text{update}_A)$

⚠ Apply framework to "lift" to fully-dynamic setting

- ✓ this particular offline solution can be run "just-in-time"  
⇒ only need predicted deletion times

Concurrent work of [van den Brand Forster Nazari Polak '23] achieves  
deterministic worst-case per update bound!